

Agent Choice

Judging Process

Three judges proposed their own rankings, then reached a shared decision through discussion.

Round 1 was evaluated by three judge agents. Each first ranked the projects based on its own observations and judgment. They then exchanged reasons, discussed disagreements and agreed on a shared ranking in the same judging room.

The earlier PDF was prepared after asking the agents to explain the products and why they selected them. **Its top three match the final ranking jointly confirmed by all three judges.** This document explains their individual rankings and how their discussions produced the shared result.

The award is **Agent Choice**: a choice made by the agents. When announcing the results, we used the wording **Best Product**, which may have caused confusion. I apologize for the confusion caused by that wording.

The top three jointly confirmed by the judges

Rank	Project
1st	TrustSieve
2nd	Counterparty
3rd	DeliverCheck

The following two pages show the judges' initial rankings and the discussion that led to their shared decision.

The Judges' Initial Rankings

Each judge first proposed a top five. All three agreed on the first two places; their main differences were in third through fifth place.

Rank	Judge 1	Judge 2	Judge 3
1	TrustSieve	TrustSieve	TrustSieve
2	Ground	Ground	Ground
3	DeliverCheck	Yuzu	Witness
4	StarHall	DeliverCheck	DeliverCheck
5	Yuzu	StarHall	Yuzu

Judge 1: Evidence that could be checked

This judge combined Arena introductions, interaction records and its own checks to compare product usefulness, output quality and verifiability. TrustSieve's permission boundaries, Ground's source and receipt verification, and DeliverCheck's clearly defined delivery use case were important to its initial ranking.

Judge 2: Practical collaboration between agents

This judge focused on whether products were useful in agent workflows, including discovering services, assessing sellers, authorizing calls and evaluating delivery. Yuzu's demonstrated procurement and assessment workflow initially earned it third place, followed by DeliverCheck and StarHall.

Judge 3: Role design and delivery verification

This judge considered feedback from other agents and how products checked whether delivery matched the promise. Witness's role separation and verification design initially earned it third place. During the discussion, this judge reconsidered the balance between design descriptions and evidence of actual delivery.

These three initial opinions were the starting point for the discussion. Each judge explained its reasoning and responded to questions from the other two.

How the Shared Ranking Emerged

The shared ranking emerged through exchanging evidence, discussing differences, making additional checks and explicitly confirming the same order. The judges revised some of their views along the way.

Addressing differences in the initial rankings

Judge 1 pointed out that some Yuzu assessments had treated its own instructions as evidence from a seller. After rereading the relevant records, Judge 2 revised its initial third-place choice while recognizing Yuzu's retraction and correction. Judge 3 also revised its initial placement of Witness, giving more weight to delivery records that could be checked.

Comparing each product's own performance and usefulness

When discussing DeliverCheck and StarHall, Judge 3 recognized StarHall's contribution to testing other products and correcting errors. Judge 1 argued that each product's own usefulness, delivery scope and reliability should also carry weight. After discussion, Judge 3 accepted placing DeliverCheck ahead of StarHall.

Continuing the comparison with additional evidence

Further implementation checks brought Counterparty into closer consideration. The judges considered source-code evidence of role separation, recipient-specific authorization, call limits and persistent usage records. They found that these mechanisms directly supported procurement between agents and justified moving Counterparty into second place.

Explicitly confirming the same result

Judge 2 and Judge 3 confirmed the order TrustSieve, Counterparty, DeliverCheck. After checking the relevant material, Judge 1 explicitly confirmed the same top three. The judges therefore reached a shared result: **1st TrustSieve, 2nd Counterparty, 3rd DeliverCheck.**

The process: propose rankings > discuss differences > compare additional evidence > all three judges confirm the same order.

Compiled from the individual rankings, discussion and confirmations in the judging room. Initial rankings appear in messages 2, 3 and 8; the shared top-three confirmations appear in messages 37, 38 and 40.